

Bi-directional GPLAR

Since inference backwards is much harder than direct learning forwards, a normal DGP running in reverse is added to be averaged with the latent function values from GPLAR. Following is a three-dimensional example: Firstly, GPLAR's kernels are defined as additive kernels between over inputs and over outputs:

$$\begin{aligned}
p(f_1|\theta_1) &= \mathcal{GP}(f_1; 0, K(x, x')) \\
p(f_2|\theta_2) &= \mathcal{GP}(f_2; 0, K(x, x') + K(f_1(x), f_1(x'))) \\
p(f_3|\theta_3) &= \mathcal{GP}(f_3; 0, K(x, x') + K([f_1(x), f_2(x, f_1(x))], [f_1(x'), f_2(x', f_1(x'))])) \\
p(g_3|\phi_3) &= \mathcal{GP}(g_3; 0, K(x, x')) \\
p(g_2|\phi_2) &= \mathcal{GP}(g_2; 0, K(g_3(x), g_3(x'))) \\
p(g_1|\phi_1) &= \mathcal{GP}(g_1; 0, K(g_2(g_3(x)), g_2(g_3(x'))))
\end{aligned}$$

Secondly, two series of latent function values are produced from the GPs.

$$\begin{aligned}
p(h_{1n}|f_1, x_n) &= \mathcal{N}(h_{1n}; f_1(x_n), \sigma^2) \\
p(h_{2n}|f_2, x_n, h_{1n}) &= \mathcal{N}(h_{2n}; f_2(x_n, h_{1n}), \sigma^2) \\
p(h_{3n}|f_2, x_n, h_{1n}, h_{2n}) &= \mathcal{N}(h_{3n}; f_3(x_n, h_{1n}, h_{2n}), \sigma^2) \\
p(h'_{3n}|g_3, x_n) &= \mathcal{N}(h'_{3n}; g_3(x_n), \sigma^2) \\
p(h'_{2n}|g_2, h'_{3n}) &= \mathcal{N}(h'_{2n}; g_2(h'_{3n}), \sigma^2) \\
p(h'_{1n}|g_1, h'_{2n}) &= \mathcal{N}(h'_{1n}; g_1(h'_{2n}), \sigma^2)
\end{aligned}$$

Lastly, at each layer/output levels, two latent function values are averaged, and final outputs are produced with noise added.

$$\begin{aligned}
p(y_{1n}|h_{1n}, h'_{1n}) &= \mathcal{N}(y_{1n}; \frac{1}{2}(h_{1n} + h'_{1n}), \sigma_1^2) \\
p(y_{2n}|h_{2n}, h'_{2n}) &= \mathcal{N}(y_{2n}; \frac{1}{2}(h_{2n} + h'_{2n}), \sigma_2^2) \\
p(y_{3n}|h_{3n}, h'_{3n}) &= \mathcal{N}(y_{3n}; \frac{1}{2}(h_{3n} + h'_{3n}), \sigma_3^2)
\end{aligned}$$

1. Put missing observations in the middle output levels.




Figure 1: Containing both close-upwards, close-downwards and middle-open observations

2. I looked into a case where predictions using close-upwards observations does not have under-estimated uncertainty. It seems like when there is no strong dependency discovered with any of the following observed outputs, (F2 in this example), GPLAR will have large uncertainty and the predicted mean will be corrected by the DGP in reverse.

Output	Time kernel variance	Output Linear kernel variances	Output Non-linear kernel variances
FZ	4.78		
F1	0.36	FZ:0.711	2e-06
F2	0.31	FZ:2.4 F1:0.21	4e-06
F3	0.60	FZ:6e-07 F1:0.63 F2:7e-07	6e-07
F4	1.90	FZ: 3.43 F1:1.27 F2:8e-06 F3:8e-05	1e-06
F5	0.28	FZ:6e-07 F1:3e-05 F2:4e-07 F3:0.51 F4:6e-07	1e-06
F6	0.59	FZ:6e-04 F1:3e-05 F2:5e-06 F3:4e-07 F4:0.19 F5:3e-04	1e-06




Figure 2: Uncertainty in F2 is not under-estimated

3. If DGP in reverse is replaced with another GPLAR in reverse, close-upwards observations can sometimes have “well-calibrated” uncertainty as shown below (This dataset is from another patient, no.345), but is not better than gpar or gpar in reverse direction in terms of log-likelihood or smse. If I remove the kernel over time for DGP in reverse, the uncertainty for first three outputs are still underestimated. And “well-calibrated” uncertainty is not always maintained, no.346 still has underestimated uncertainty in first two outputs, while no.347 has underestimated uncertainty in the last two outputs.




Figure 3: Patient No.345 using Bi-GPLAR




Figure 4: Patient No.345 with shorter missing area using Bi-GPLAR




Figure 5: GPLAR + DGP in reverse with time kernel removed

output	log-likelihood			smse		
	gplar	gpar	gpar-r	gplar	gpar	gpar-r
FZ	-622.1	-59.69	-145.3	0.2473	0.4062	0.5754
F1	-196.4	-60.33	-34.98	0.0597	0.3949	0.0422
F2	-42.67	-76.34	-37.17	0.0297	0.6600	0.1322
F3	-6.245	-20.43	-55.52	0.0293	0.0233	0.4367
F4	-86.05	-25.43	-62.93	0.0529	0.0779	0.5360
F5	-41.07	-60.58	-50.71	0.1605	0.2360	0.7496
F6	-63.48	-32.88	-117.8	0.0440	0.0192	0.4358




Figure 6: Patient No.346




Figure 7: Patient No.347




Figure 8: I tried smaller missing area on exchange-rate dataset using GPLAR in both direction




Figure 9: Some more experiment with close-upwards observations




Figure 10: Some interesting behaviour on 2006 datasets




Figure 11: Some interesting behaviour on 2000 datasets

It is possible that GPAR will be over-fitted to outliers, however, GPLAR might also underfit as shown in Figure 11. We can observe in this exchange-rate dataset that GPAR tends to fit data by more fluctuations, smaller length-scales in kernels, while GPLAR will cover data by increasing the observation noise.

In multi-fidelity model, they conduct an experiments such that 2005's data is considered as low-fidelity input while 2015's data is considered as high-fidelity. I did the same with exchange-rate datasets, collecting USD/GBP data from year 2000-2009, and try to predict some missing values in 2009. For every year, the input "time" will be different. However, it seems like the dependencies between years inside same group are rather weak, as shown from the learnt variances value of kernels.




Figure 12: Some interesting behaviour on 2000 datasets